

Lecture 8

Memory & Cache

Overview

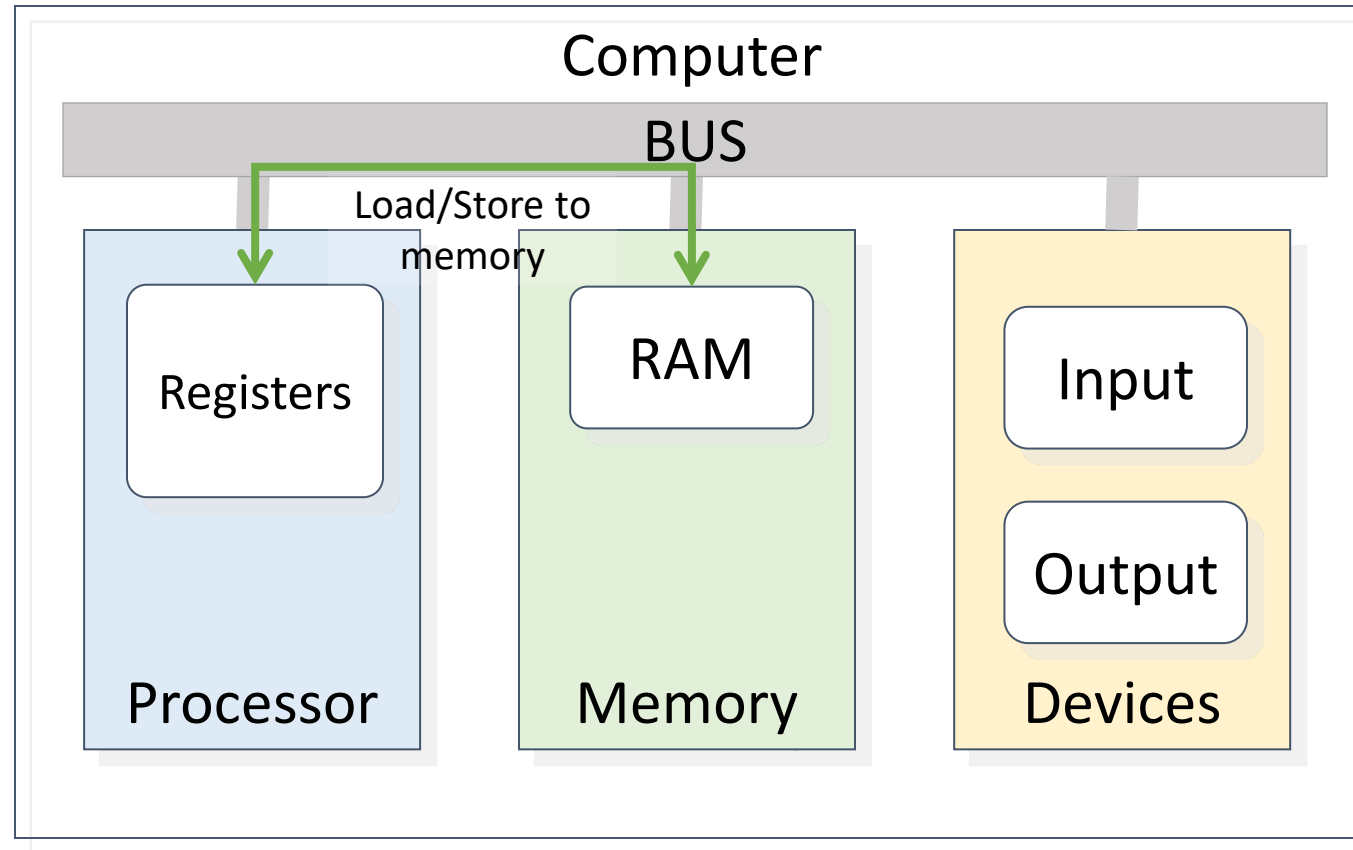
Memory

- Program usage of memory
- Hardware
 - Volatile vs Non-Volatile Memory
 - DRAM vs SRAM
 - Current trend of DRAM

Cache

- What is it?
- How does it work?
- Current trend of Cache

Data Transfer: The BIG picture



Registers are fast storage in the processor. If operands (data values) are in memory we have to **load** them to processor (registers), operate on them, and **store** them back to memory

Main Memory Storage

Physical memory storage

- Random Access Memory (RAM)
- Can be viewed as an array of bytes
- Each byte has a unique index
 - Known as [physical address](#)

content	address
⋮	
	1024
	1025
	1026
	1027
	1028
⋮	


Variable and Pointer

```
int x;  
int *ptr;  
ptr = &x;  
*ptr = 123;
```

Normal variable is referred by variable name in code

– No need to know the address

A pointer contains the address of a memory location



Pointers and Arrays


Array name is a **constant pointer**

- Points to the zeroth element

```
int ia[3] = {3, 5, 7};
```

- Is the following valid?

```
int* ptr;  
  
ptr = ia;  
ia = ptr;  
ptr[2] = 9;  
  
ptr = &ia[1];  
ptr[1] = 11;
```



Source Code → Executable

Executable typically contains:

- Code and Data layout

```
int i = 0;

void inc()
{
    i = i + 20;
}

int main()
{
    inc();
    ...
    ...
}
```

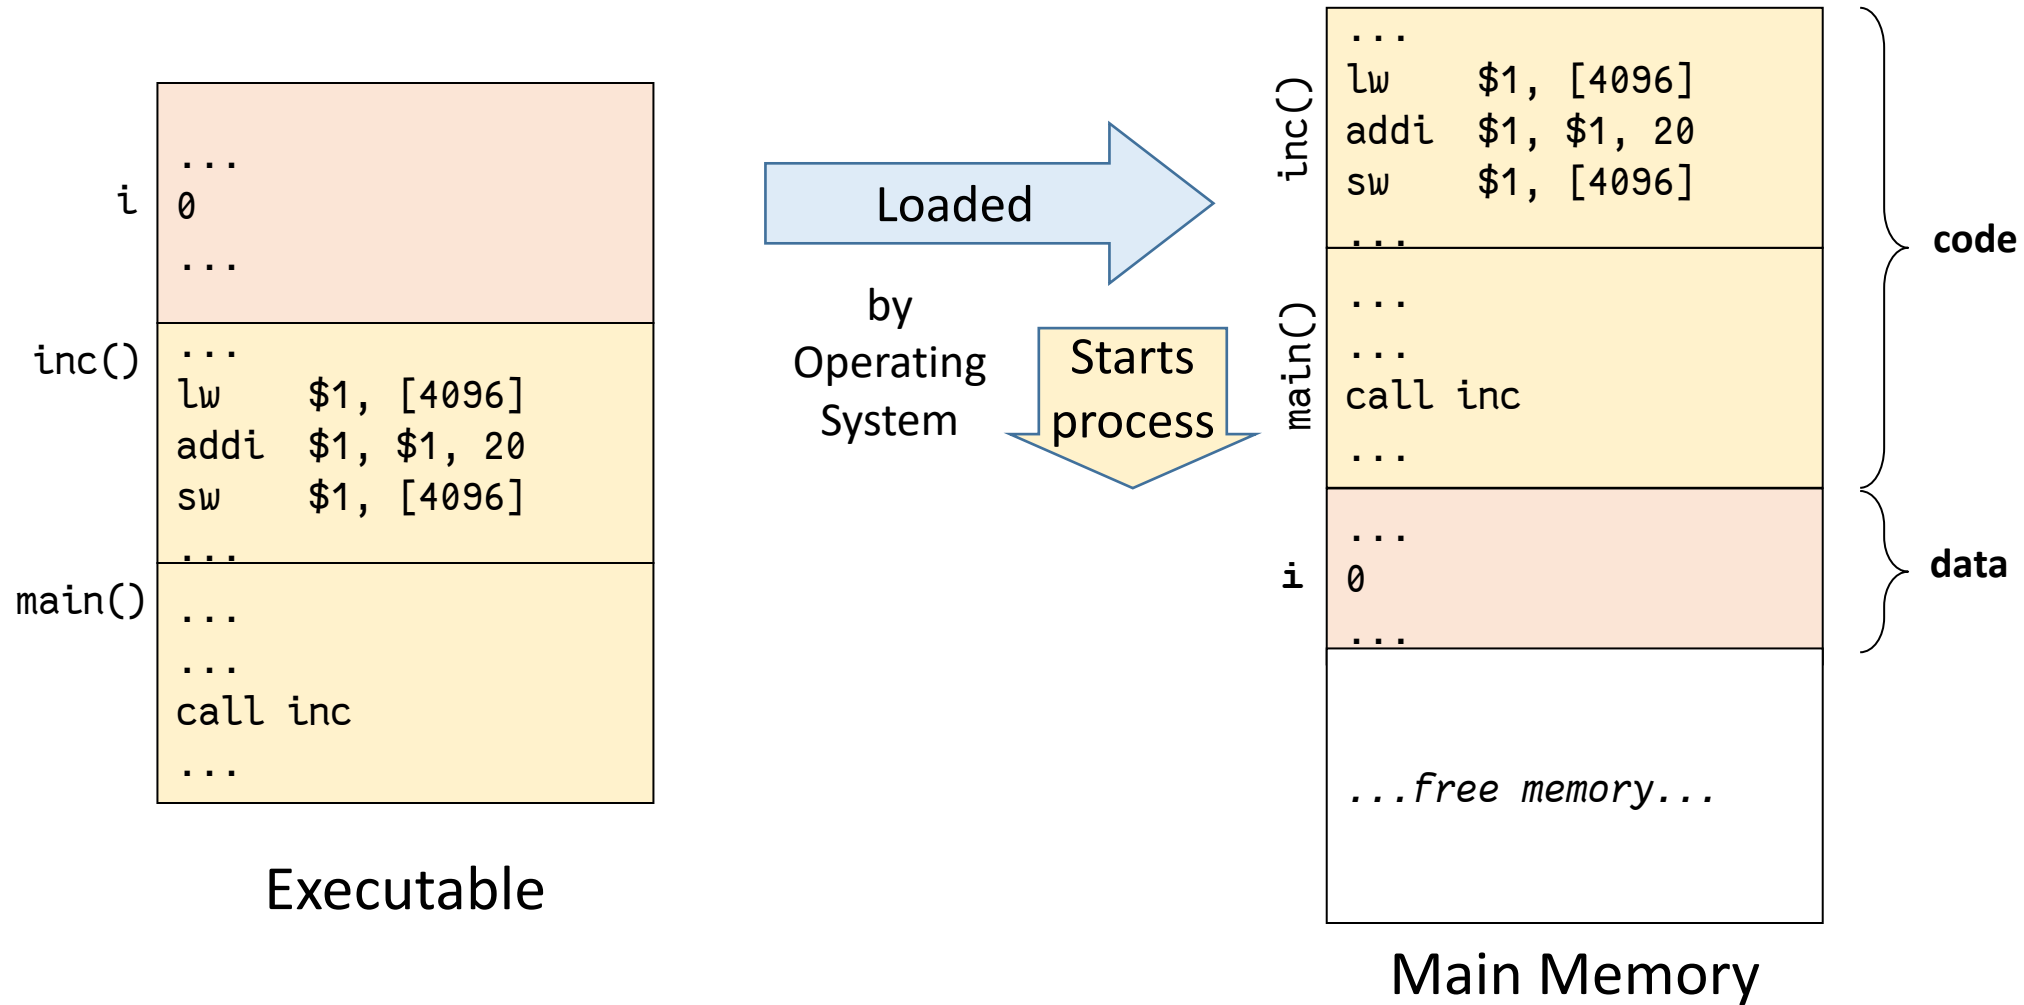
Source Code

Compile


i	...
	0
	...
inc()	...
	lw \$1, [4096]
	addi \$1, \$1, 20
	sw \$1, [4096]
	...
main()	...
	...
	call inc
	...

Executable

Executable → Loaded into Memory



Memory Usage of Process Simplified



How about variables defined in a function?

- e.g. Parameters, Local variables

Variables in function: Scoping

```
void f()
{
    int x;
    g(123);
    ...
}
void g(int x)
{
    h();
    ...
}
void h()
{
    int x;
}
```

- Consider the example of the left
- The 3 "x" variables are actually different from each other
- Also, these "x" exists **only when** the respective function is executing!

Stack Memory

Stack Memory Region


- New memory region (environment) to store data for function calls
- Environment of each function call is known as a stack frame

Stack frame contains:

- Arguments (actual parameters) for the function
- Local variables
- Other information.... (not covered in this course)

Illustration: Stack Memory Usage (1 / 5)

```
void f()  
{  
    int x;  
    g(123);  
    ...  
}  
  
void g(int x)  
{  
    h();  
    ...  
}  
  
void h()  
{  
    int x;  
}
```



At this
point

Stack Frame
for f() **x**





Illustration: Stack Memory Usage (2 / 5)

```
void f()  
{  
    int x;  
    g(123);  
    ...  
}
```

```
void g(int x)  
{  
    h();  
    ...  
}
```

```
void h()  
{  
    int x;  
}
```



At this
point

Stack Frame
for g()

x

Stack Frame
for f()

x





Illustration: Stack Memory Usage (3 / 5)

```
void f()  
{  
    int x;  
    g(123);  
    ...  
}  
  
void g(int x)  
{  
    h();  
    ...  
}  
  
void h()  
{  
    int x;  
}
```



At this
point

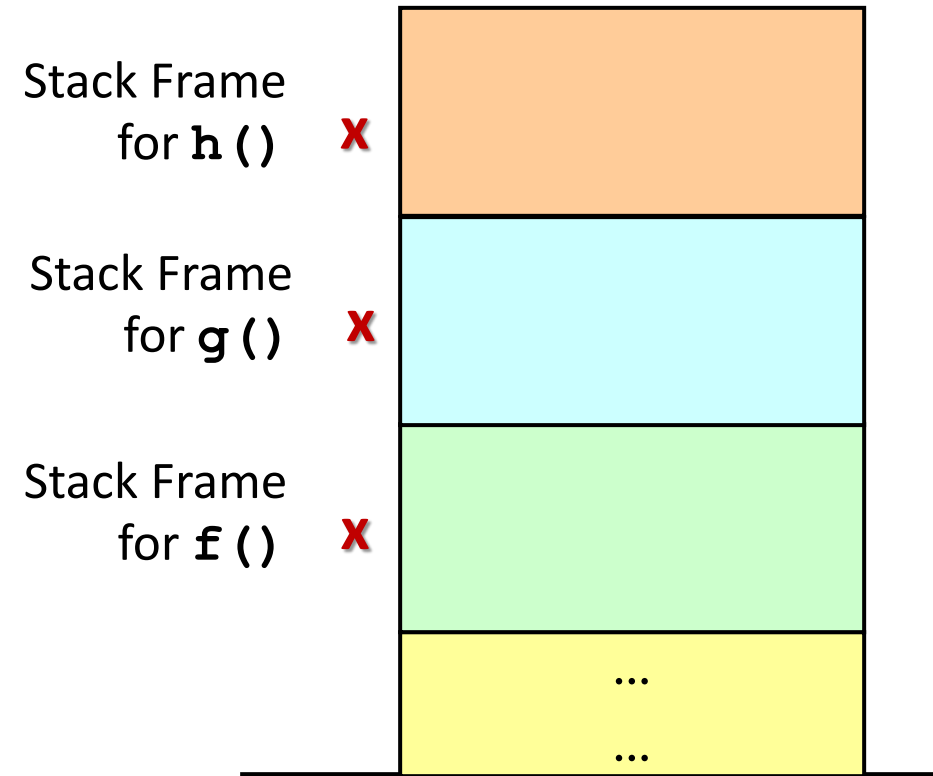



Illustration: Stack Memory Usage (4 / 5)

```
void f()  
{  
    int x;  
    g(123);  
    ...  
}
```

```
void g(int x)  
{  
    h();  
    ...  
}
```



At this
point

```
void h()  
{  
    int x;  
}
```

Stack Frame
for g()

x

Stack Frame
for f()

x





Illustration: Stack Memory Usage (5 / 5)

```
void f()  
{  
    int x;  
    g(123);  
    ...  
}
```




At this
point

```
void g(int x)  
{  
    h();  
    ...  
}
```

```
void h()  
{  
    int x;  
}
```

Stack Frame
for f() **x**



Stack Memory: Key Ideas

Stack frames are created / destroyed at runtime:

- Entering a function: new stack frame
- Exiting a function: stack frame destroyed

Hence, we can have multiple variables with the same name in different function (and scope)

Function : Pass by value


```
void swap_ByValue( int a, int b )
{
    int temp;

    temp = a;
    a = b;
    b = temp;
}

int main()
{
    int i = 123, j = 456;

    swap_ByValue( i, j );

    cout << i << endl;
    cout << j << endl;
}
```



Function : Pass by address/pointer


```
void swap_ByAdr( int* a, int* b )
{
    int temp;

    temp = *a;
    *a = *b;
    *b = temp;
}

int main()
{
    int i = 123, j = 456;

    swap_ByAdr( &i, &j );

    cout << i << endl;
    cout << j << endl;
}
```



Function : Pass by reference


```
void swap_ByRef( int& a, int& b )
{
    int temp;

    temp = a;
    a = b;
    b = temp;
}

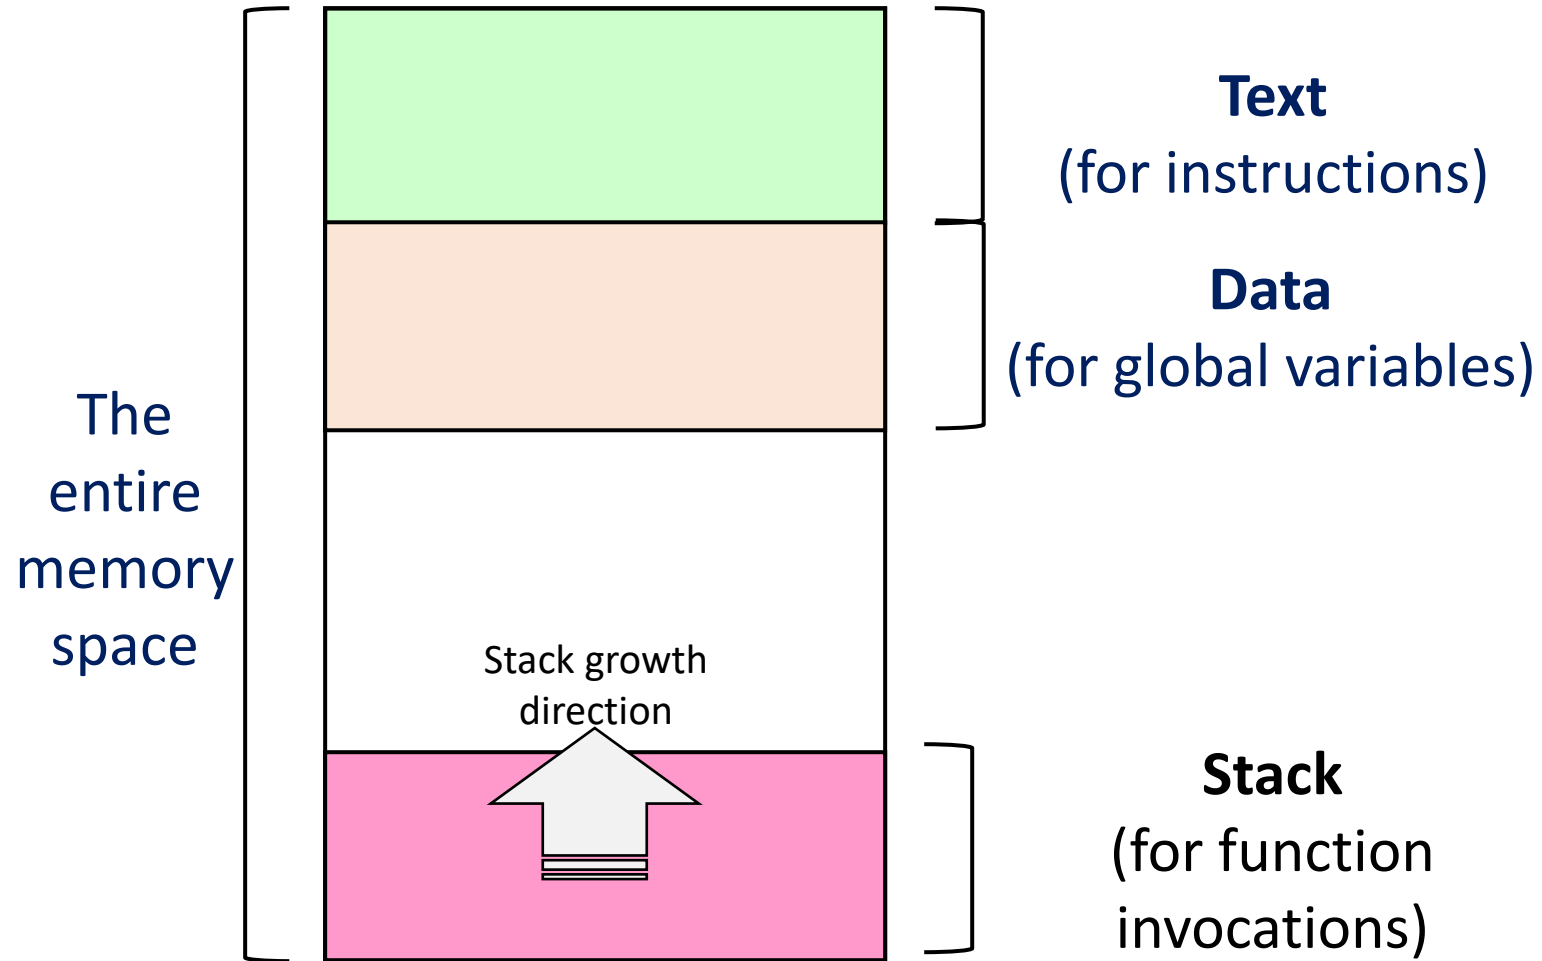
int main()
{
    int i = 123, j = 456;

    swap_ByRef( i, j );

    cout << i << endl;
    cout << j << endl;
}
```



Memory Usage of Process Updated



Memory Hardware

Volatile and Non-Volatile Memory

Volatile Memory (main focus):

- Storage that loses information when electrical power is interrupted
- e.g. RAM (Random Access Memory)
- Main storage of Process Memory discussed

Non-Volatile Memory

- Storage that can maintain information even without power
- e.g. ROM (Read-only memory), Harddisk, SSD (Solid State Drive), etc.
- Commonly store only persistent information, e.g. file

Memory Technology : 1950s



1948: Maurice Wilkes examining EDSAC's delay line memory tubes
16-tubes each storing 32 17-bit words



Maurice Wilkes: 2005





Figure 10. IBM 2361 Core Storage

1952: IBM 2361 16KB magnetic core memory

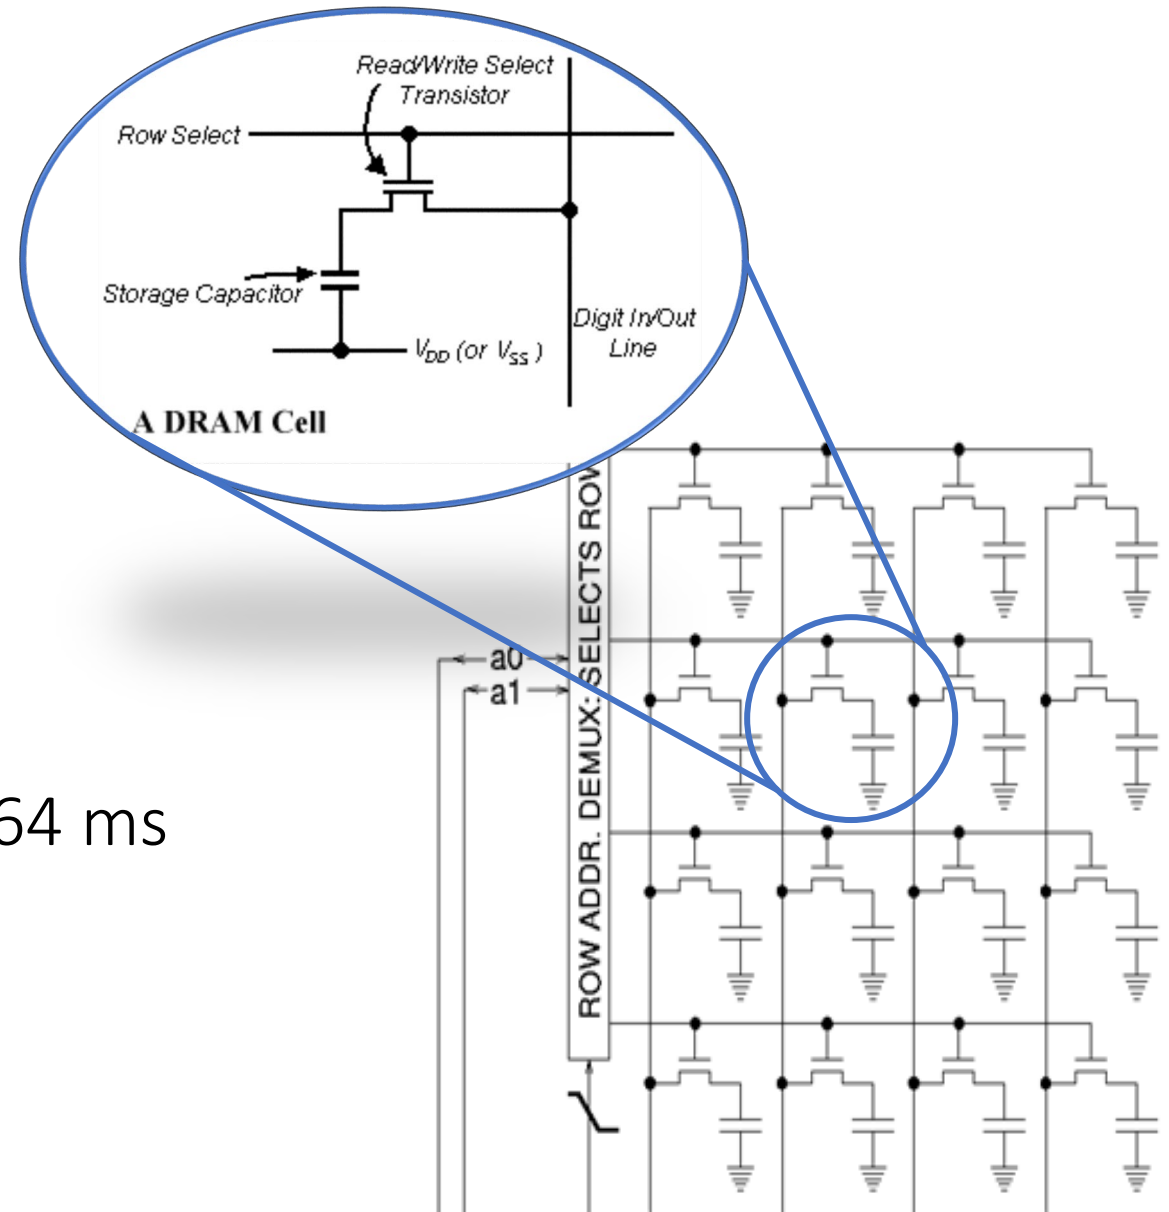
Dynamic RAM

High density

- 1 transistor + 1 capacitor per memory cell
- Cheap!

Slow access

- Latency of 50-70 ns
- Require periodic refresh every 64 ms
 - Or each cell every 7.8 μ s
- Memory blocked for 75-120 ns




16 DRAM cells, i.e. 16bits

Memory Technology Today: SDRAM


DDR SDRAM

- Double Data Rate Synchronous Dynamic RAM
- The dominant memory technology in PC market
- Delivers memory on the positive and negative edge of a clock (double rate)



DDR SDRAM Standard	Internal rate (MHz)	Bus clock (MHz)	<u>Prefetch</u>	Data rate (MT/s)	Transfer rate (GB/s)	Voltage (V)
SDRAM	100-166	100-166	1n	100-166	0.8-1.3	3.3
DDR	133-200	133-200	2n	266-400	2.1-3.2	2.5/2.6
DDR2	133-200	266-400	4n	533-800	4.2-6.4	1.8
DDR3	133-200	533-800	8n	1066-1600	8.5-14.9	1.35/1.5
DDR4	133-200	1066-1600	8n	2133-3200	17-21.3	1.2


Snapshot: DDR RAM



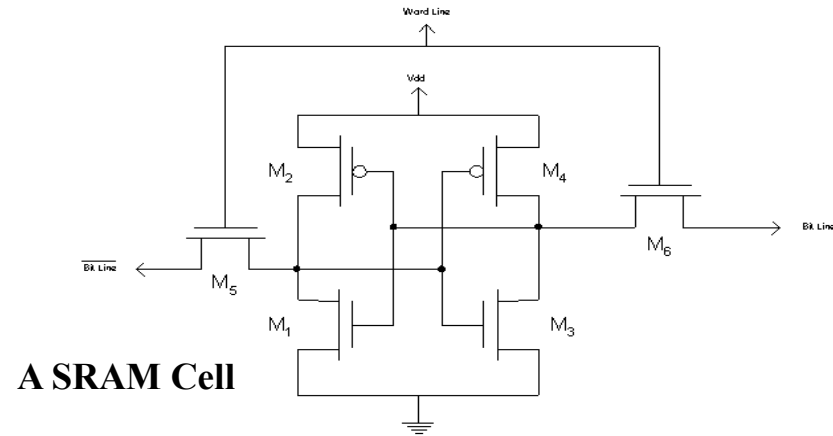
Processor-DRAM Performance Gap

Memory Wall:


- 1 GHz Processor \rightarrow 1 ns per clock cycle
- 50 ns for DRAM access \rightarrow 50 processor cycles per memory access!!



Faster Memory Technology: Static RAM



A SRAM Cell (Alternative Drawing)

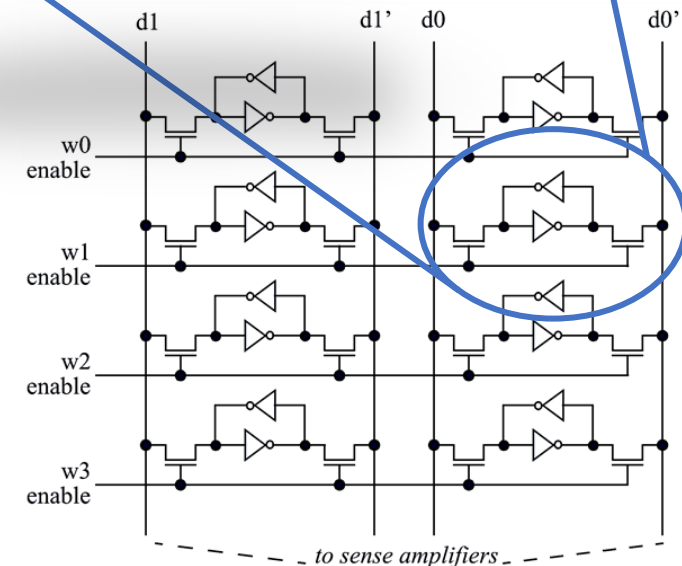


Low density

- 6 transistors (2 inverters) per memory cell
- Expensive

Fast access

- Latency of 0.5-5 ns
- No refresh needed



8 SRAM Cells, i.e 8 bits


Slow Memory Technologies: Magnetic Disk

Typical high-end hard disk:

- Average Latency: 4 - 10 ms
- Capacity: 500-5,000GB



Drive Physical and Logical Organization



Slow Memory Technologies: Flash Memory

Used in


- Thumbdrives
- Memory Cards
- Solid State Drives (SSD)

Specially designed transistors

- Trap electrons between gates


High-end SSD

- Latency of 50 μ s



Credit: Original picture from <http://encyclopedia2.thefreedictionary.com/EEPROM>

Quality vs Quantity



	Capacity	Latency	Cost/GB
Register	100s Bytes	20 ps	\$\$\$\$
SRAM	100s KB	0.5-5 ns	\$400/GB
DRAM	10s GB	50-70 ns	\$8/GB
Hard Disk	1000s GB	5-20 ms	\$0.025/GB
Ideal	1 TB	1 ns	Cheap

Best of Both Worlds

What we want

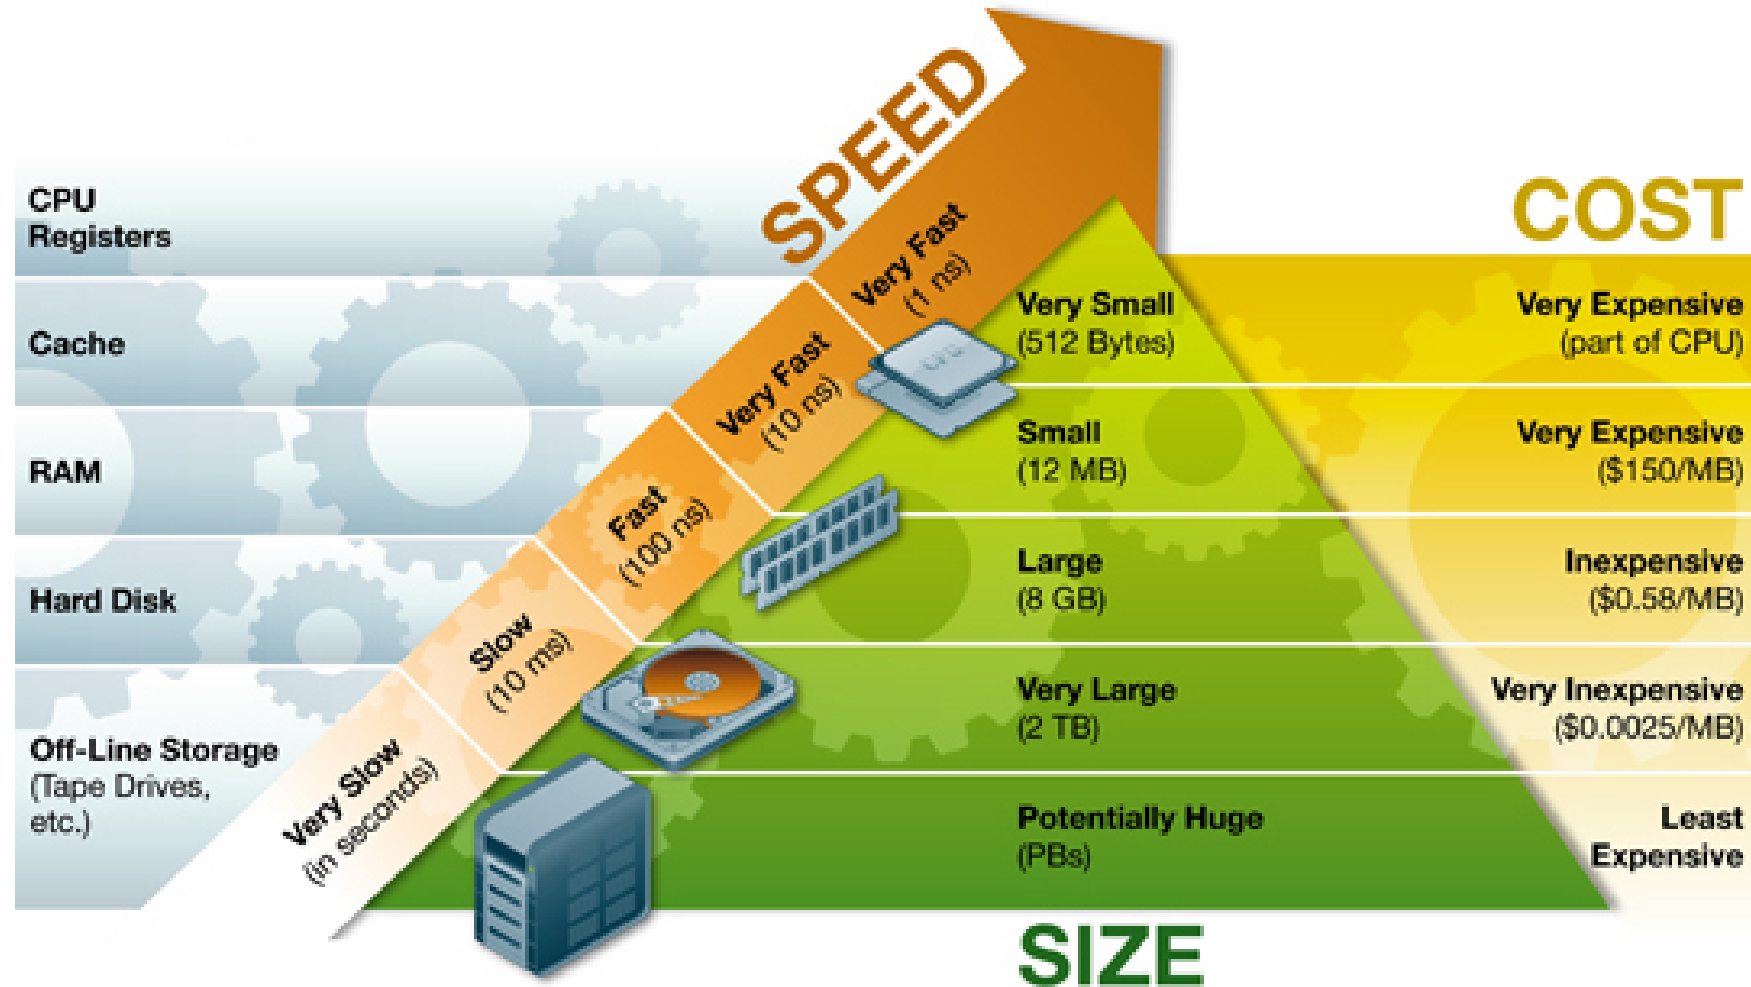
- A BIG and FAST memory
- Memory system should perform like 1GB of SRAM (1 ns access time) but cost like 1GB of slow memory

Key Idea

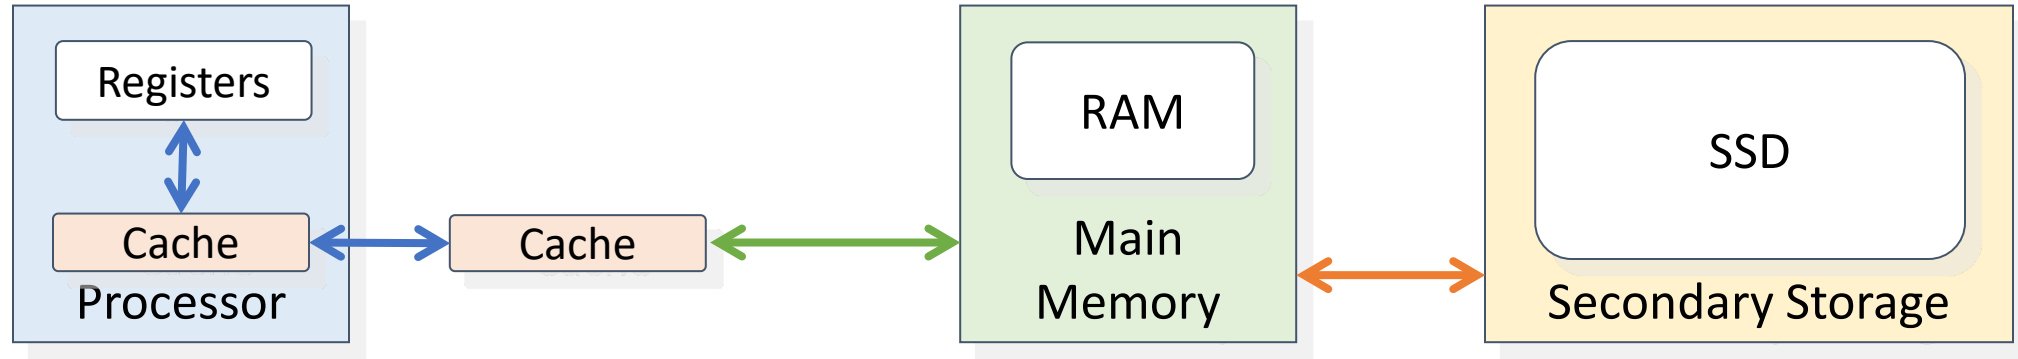
Use a hierarchy of memory technologies

- Small but fast memory near CPU
- Large but slow memory farther away from CPU

The Memory Hierarchy



Two Aspects of Memory Access



make **SLOW** main memory appear **faster**?

- **Cache**: a small but fast SRAM near the CPU
- Hardware managed: Transparent to programmer

make **SMALL** main memory appear **bigger**?

- **Virtual memory**
- OS managed: Transparent to programmer

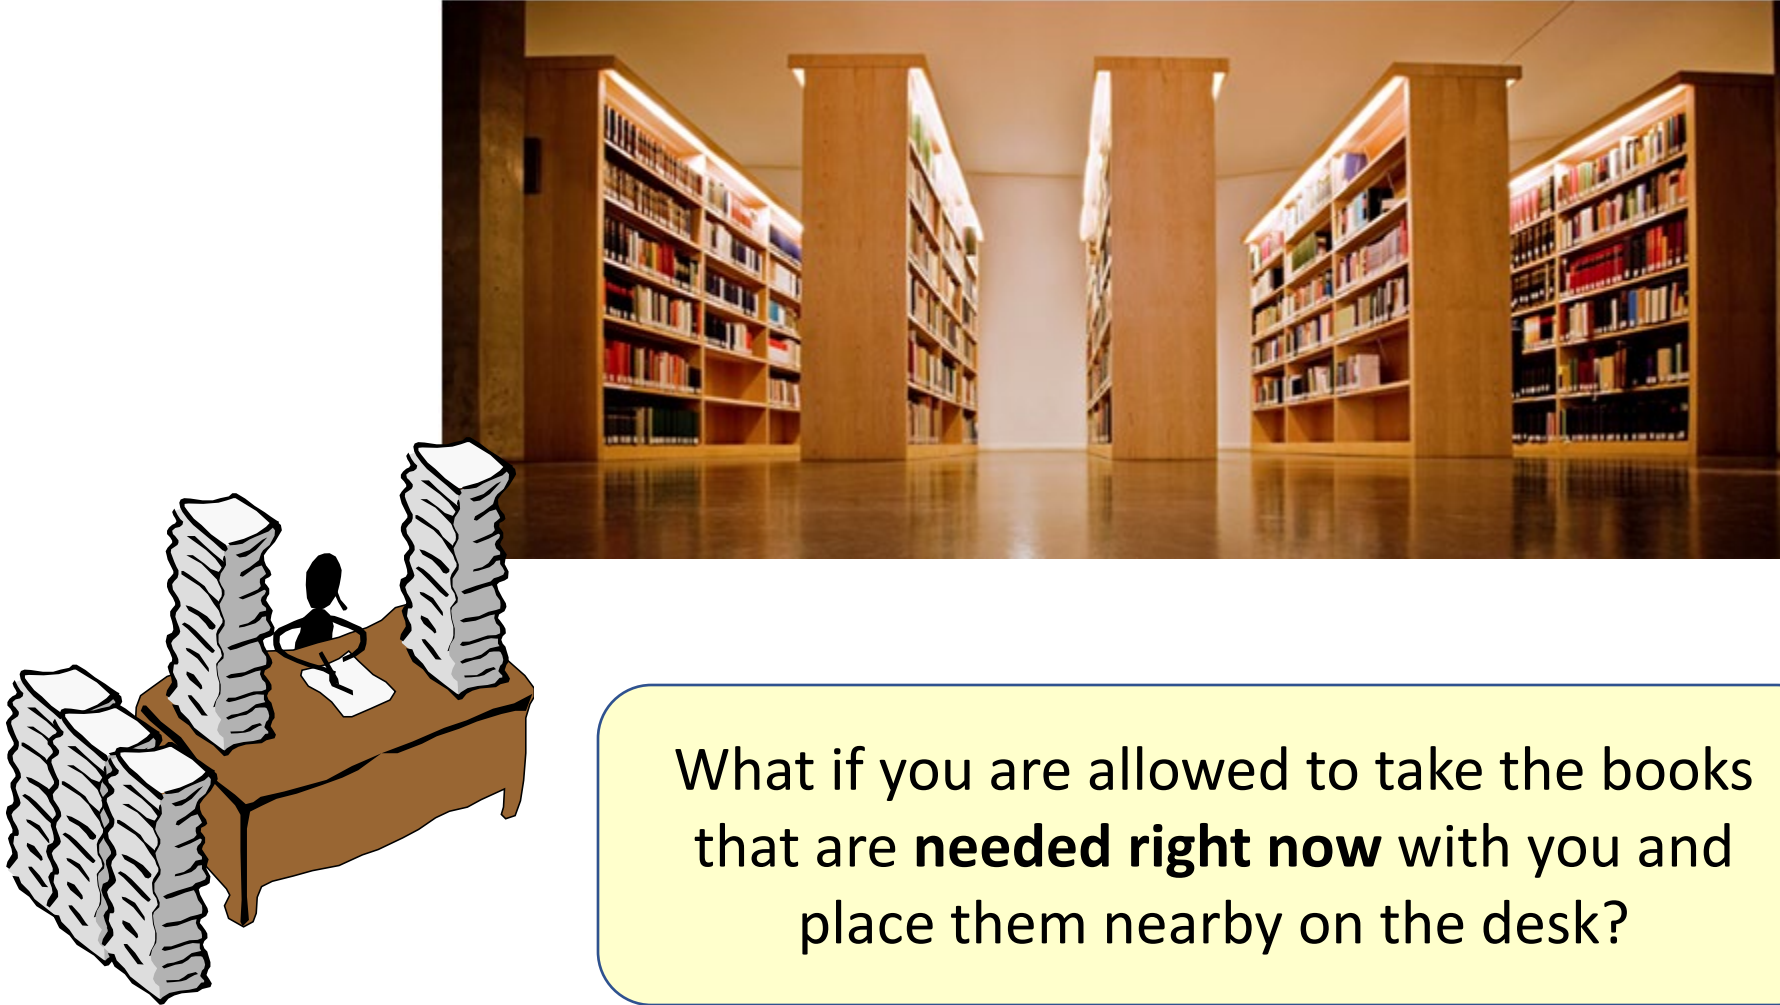
Cache

The Library Analogy



Imagine you are forced to put back a book to its bookshelf before taking another book.....

Solution: Book on the Table!



What if you are allowed to take the books that are **needed right now** with you and place them nearby on the desk?

Cache: The Basic Idea

Keep the frequently and recently used data

- in smaller but faster memory

Refer to bigger and slower memory:

- Only when you cannot find data/instruction in the faster memory

Why does it work?

Principle of Locality

Program accesses only a small portion of the memory address space within a small time interval

Types of Locality

```
void cumul_sum(int arr[], int size) {  
    for (int i = 1; i < size; ++i) {  
        arr[i] = arr[i] + arr[i-1];  
    }  
}  
  
int main(void) {  
    int i, sq[SIZE];  
  
    for (i = 0; i < SIZE; i++) {  
        sq[i] = i* i;  
    }  
  
    cumul_sum(sq, 10);  
    return 0;  
}
```

- **Temporal locality**

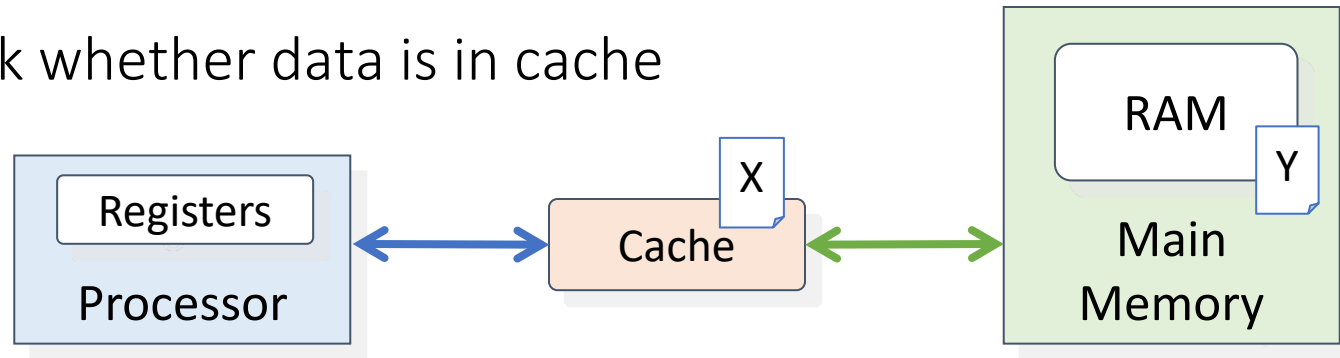
- If an item is referenced, it will tend to be referenced again soon

- **Spatial locality**

- If an item is referenced, nearby items will tend to be referenced soon

Cache: Basic Usage and Terminology

Processor check whether data is in cache



Cache Hit: Data is in cache (e.g., X)

- Data is returned to processor from cache
- This is the ideal case as cache is fast

Cache Miss: Data is not in cache (e.g., Y)

- Data is loaded from memory
- [Simplified] Data is placed in cache for future reference
- Will incur the full memory latency (i.e. slow)

MEMORY \leftrightarrow CACHE MAPPING

Where to place the memory block in cache?

Terminology: Cache Block/Line


Cache Block/Line:

- Unit of transfer between memory and cache

Block size is typically one or more words

- e.g.: 16-byte block \cong 4-word block
- 32-byte block \cong 8-word block

Why block size is bigger than word size?



Fully Associative Cache

Having total freedom

Fully-Associative (FA) Analogy

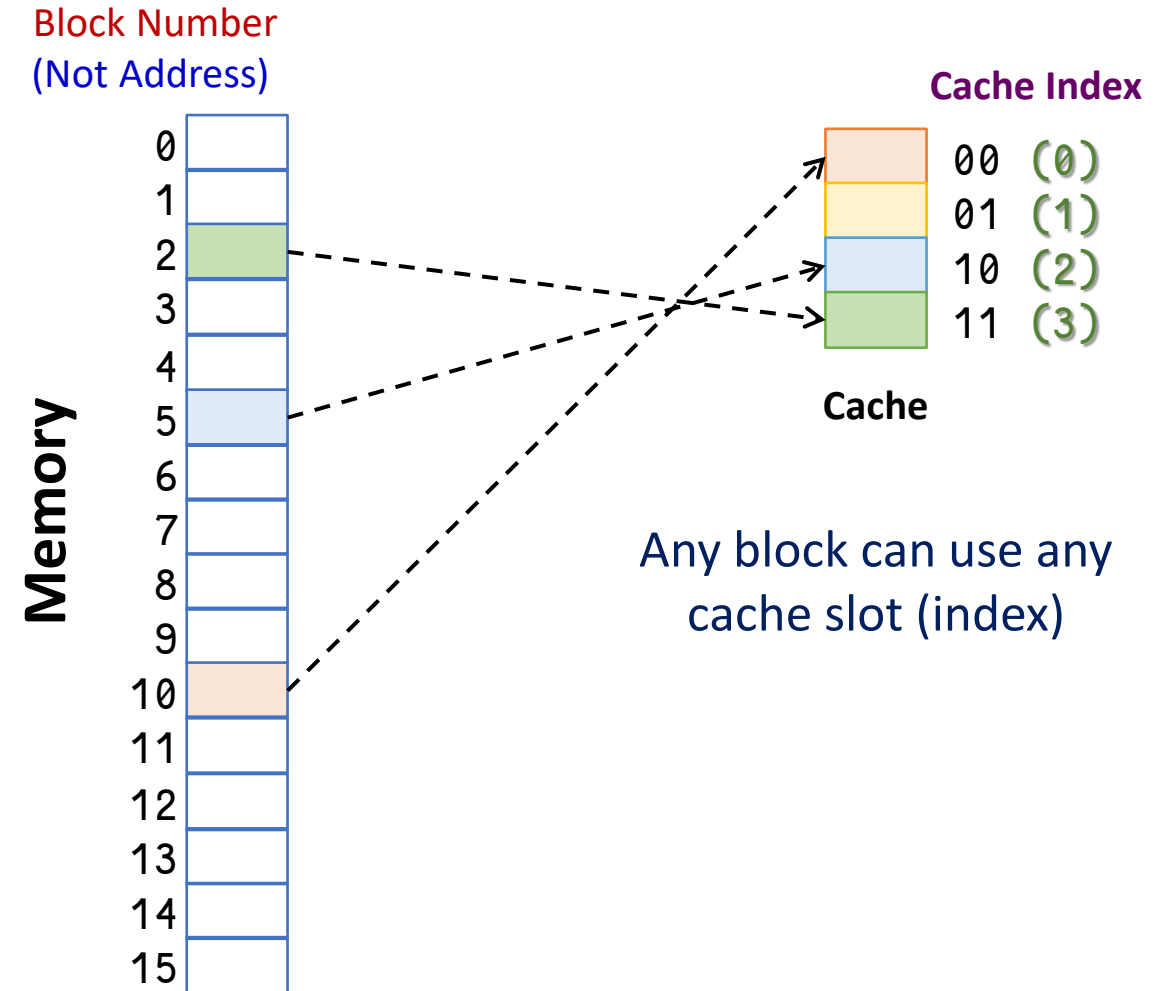


A book can go into any location on the desk!

Fully Associative Cache

Question:

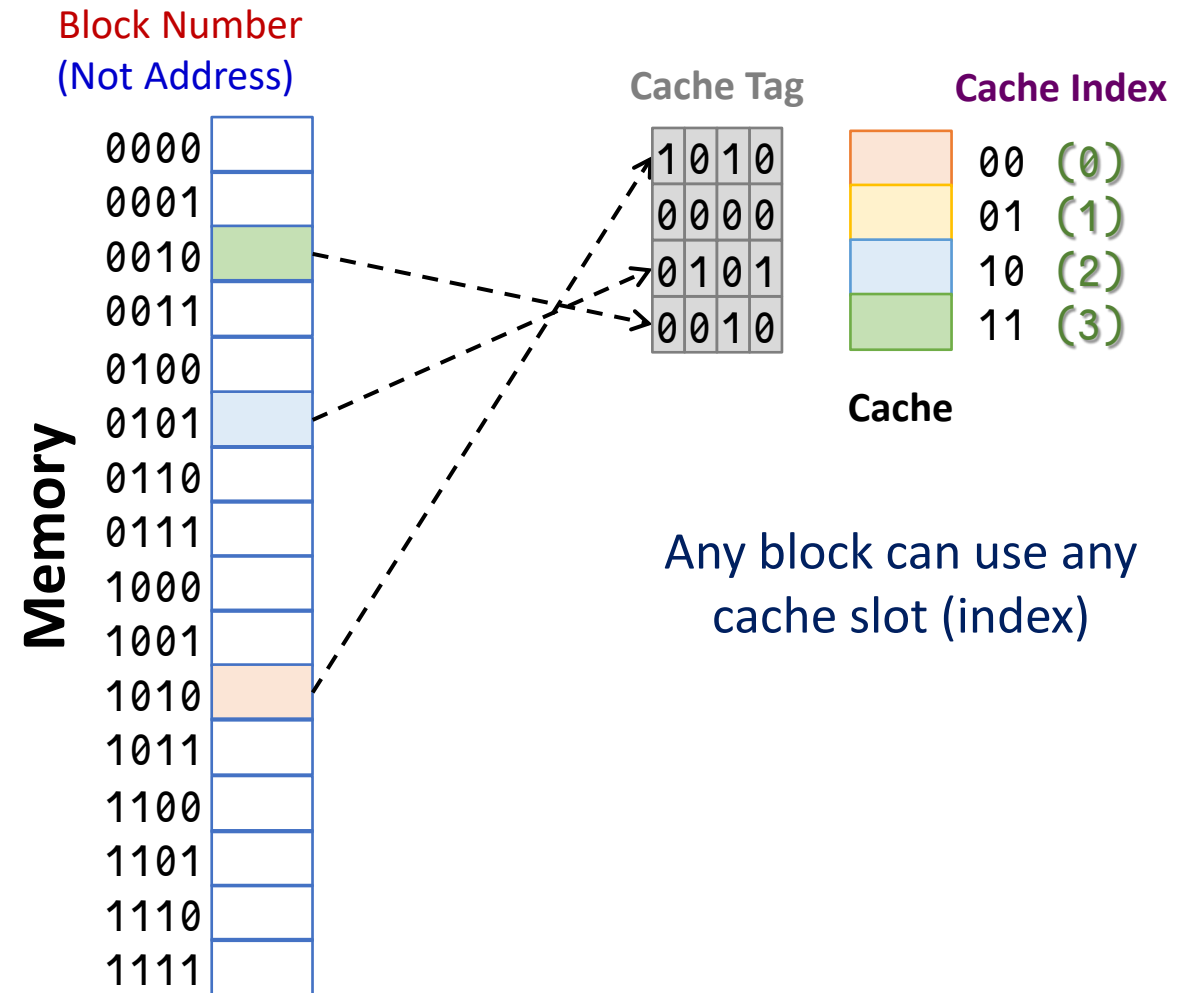
- How to identify which block is in the cache?



Fully Associative Cache

Cache Tag

- An identifier (id) to identify the current block in the cache
- Entire block number must be used



FA Cache: Pros / Cons

Intuitive

- Any block can go into any cache location

Resource intensive

- Need to search all locations at once
- need a lot of hardware comparators
- take up significant space on chip!

➔ Fully associative cache is seldom used

➔ If used, usually in very limited capacity

Direct Mapped Cache

Let's restrict the freedom a little?

Direct Mapping (DM) Analogy

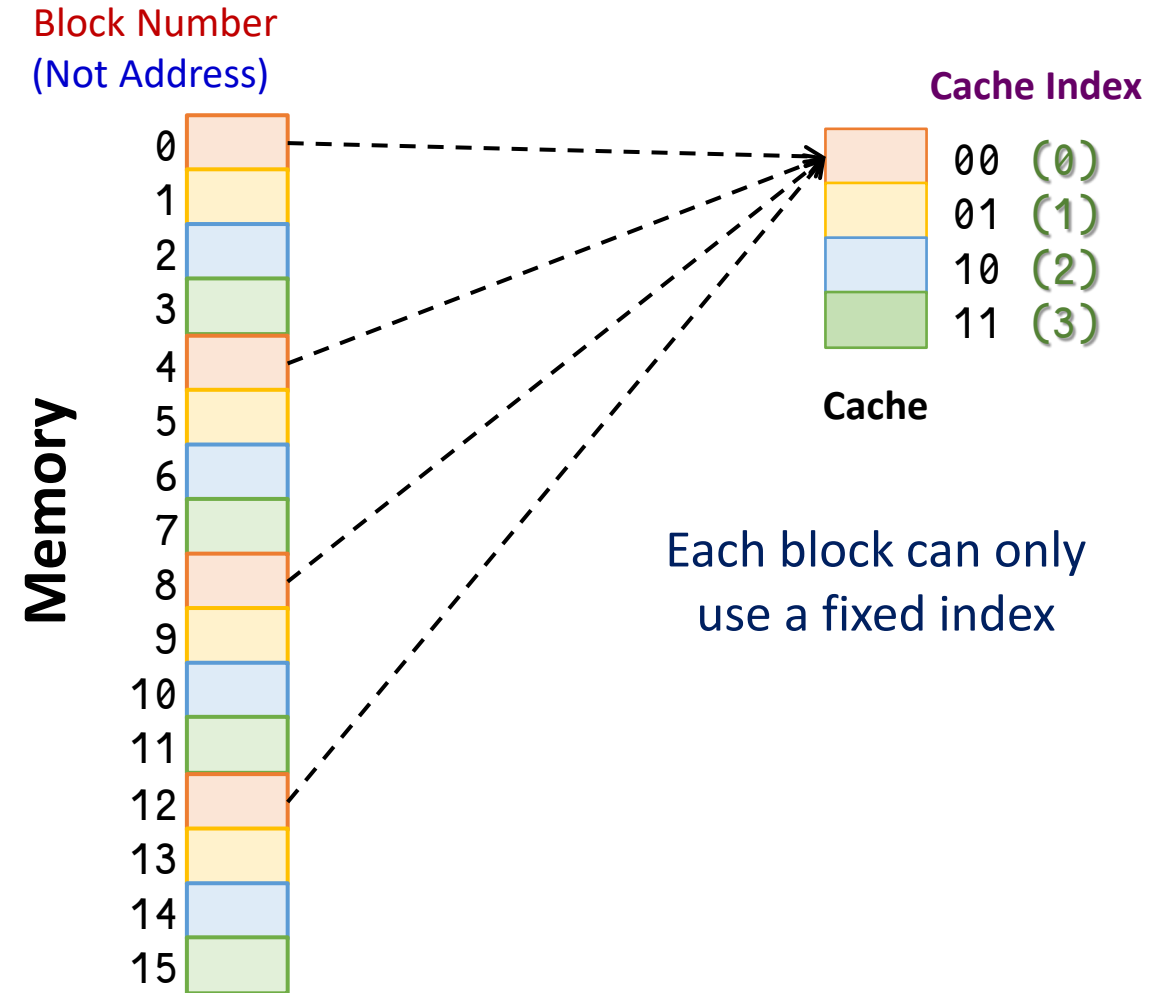


Imagine there are 26 "locations" on the desk to store book. A book's location is determined by the first letter of its title.
➔ Each book **has exactly one location**

Direct Mapped Cache

Question:

- Given a block number, identify the cache index it should go



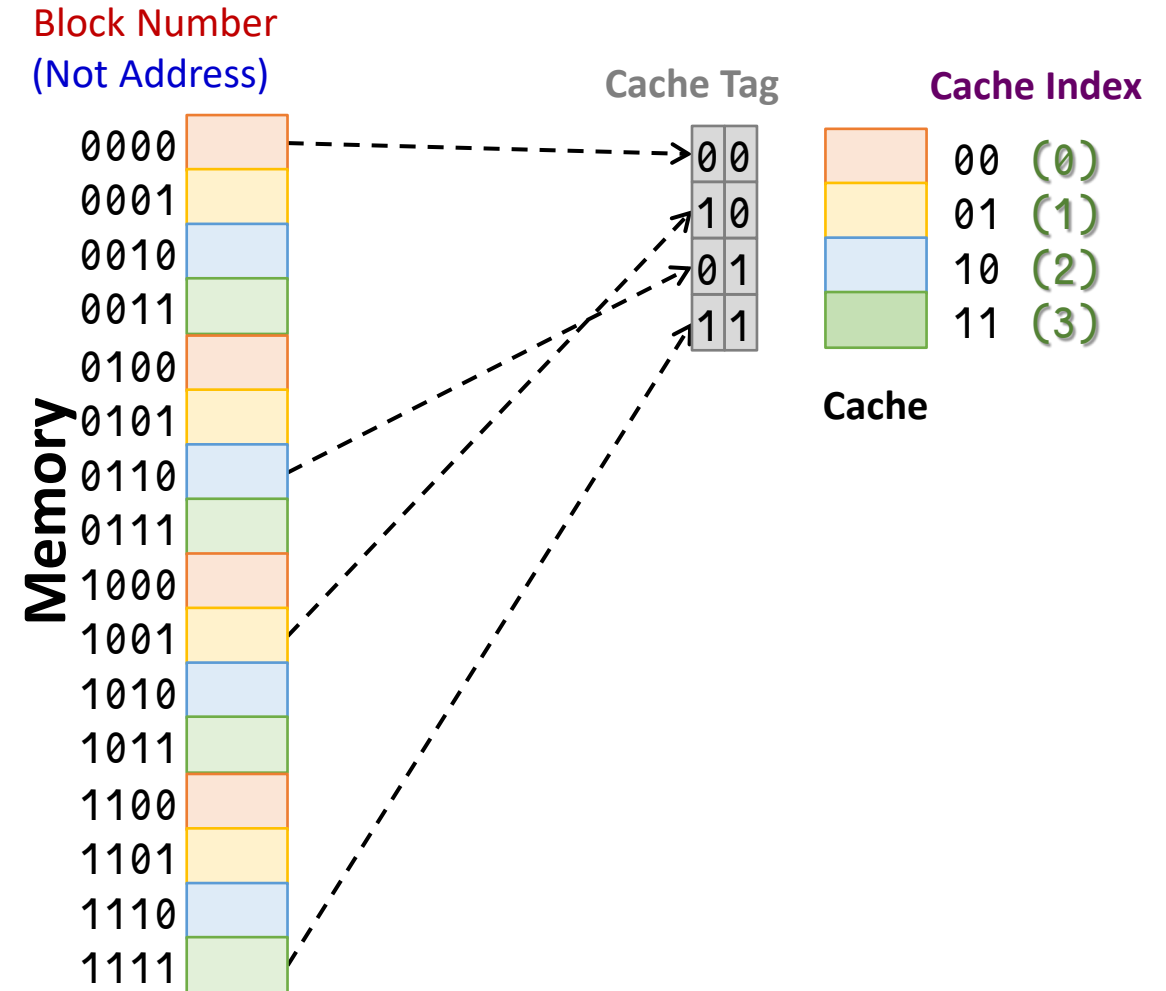
Direct Mapped Cache

Question:

- Given a block number, identify the cache index it should go
- How to identify block in cache?

Cache Tag

- Only requires most significant part of block number



DM Cache Pros / Cons

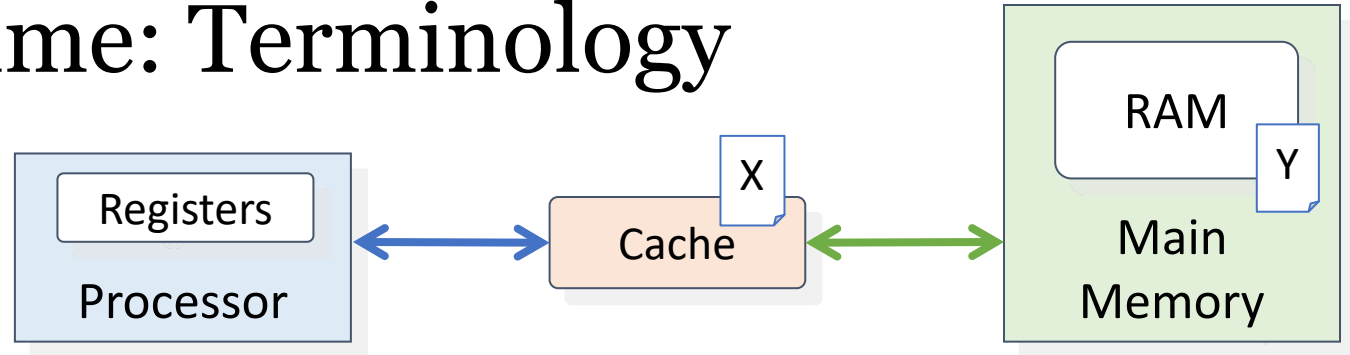
Better use of hardware resource:

- Only one possible cache location for any block
- Can have 1 single hardware comparator

Potential of cache conflicts:

- What if we happen to use different blocks that maps to the same cache index?

Memory Access Time: Terminology



Hit: Data is in cache (e.g., X)

- Hit rate: Fraction of memory accesses that hit
- Hit time: Time to access cache

Miss: Data is not in cache (e.g., Y)

- Miss rate = $1 - \text{Hit rate}$
- Miss penalty: Time to replace cache block + deliver data

Hit time < Miss penalty

Memory Access Time: Formula

$$\text{Average Access Time} \\ = \text{Hit rate} \times \text{Hit Time} + (1 - \text{Hit rate}) \times \text{Miss penalty}$$

Example:

Suppose our on-chip SRAM (cache) has **0.8 ns** access time, but the fastest DRAM we can get has an access time of **10ns**. **How high a hit rate** do we need to sustain an average access time of **1ns**? For simplicity, we assume the miss penalty is the same as the DRAM access time.

Cache: Exploration

Alternative mapping scheme:

- Set associative (hybrid between FA and DM)

More block replacement algorithm:


- Which block to replace when the cache is full?

Multi-level cache

We simplified:

- Cache address calculation
- Memory access between a load / store

Cache Sample: Intel Core i7-3960K



-2.27 billion transistors

-15MB shared Inst/Data Cache (LLC)

Per Core:

-32KB L1 Inst Cache

-32KB L1 Data Cache

-256KB L2 Inst/Data Cache

-up to 2.5MB LLC

END